

# Akhil Theerthala

Bengaluru | akhiltvsn@gmail.com | 8179528501 | LinkedIn | Portfolio | HuggingFace

## Summary

Applied Scientist with current research interests in advancing domain-specific reasoning and interactive decision-support systems. Open-source contributor with **33,000+ model family downloads** (Kuvera) and a **Global Reasoning Challenge Winner**. Proven industry impact: achieved **97.5% latency reduction** in document processing and **27.6% accuracy boost** in table detection.

## Professional Experience

<b>Senior Member Data Scientist</b>   <i>Perfios Software Solutions</i>	Apr 2025 – Present
<ul style="list-style-type: none"><li>Established a new baseline TEDS score of <b>0.85</b> on complex financial layouts by adapting open-weight VLMs (PaliGemma2) via Low-Rank Adaptation (LoRA) on domain-specific datasets.</li><li>Architected multi-stage agentic reasoning workflows using Google-ADK for underwriting, customer retention risk and claim summarization for new product lines.</li><li>Developed a Reference-Free algorithm to quantify document legibility, utilizing a regressive Vision Transformer (ViT). Filtered low-fidelity inputs with <b>92% precision</b>, preventing downstream hallucination and reducing compute costs.</li></ul>	
<b>Member Data Scientist</b>   <i>Perfios Software Solutions</i>	Jun 2023 – Apr 2025
<ul style="list-style-type: none"><li>Resolved inference latency bottlenecks (<math>8s \rightarrow 200ms</math>) by compressing a large multimodal classifier into a lightweight student architecture via Distillation and quantization techniques, preserving statistical parity in F1-scores.</li><li>Boosted generalized table detection accuracy by <b>27.6%</b> through the careful, semi-synthetic data curation for finetuning model and rigorous evaluation of YOLOv8 models.</li><li>Augmented vision-only Table Structure Recognition (TSR) systems with a semantic row-detection module. Fine-tuned specialized text-encoders to capture lexical dependencies in bank-statements, with real-time inference (&lt;40ms).</li></ul>	

## Additional Experience

<b>Research Volunteer</b>   <i>Financial Services Innovation Lab, GeorgiaTech (External Collaborator)</i>	Aug 2025 – Present
<ul style="list-style-type: none"><li><b>AAAI 2026 Acceptance:</b> Co-authored FinForge, a semi-synthetic benchmark generation pipeline for financial tasks. (Accepted at AAAI Workshop on Agentic AI).</li><li><b>Resaerch Focus:</b> Contributing to the <b>FinGT (Financial Generative Transformers)</b> project, focusing on enhancing financial domain adaptation and logic adherence in LLMs.</li></ul>	
<b>Open Source Contributor</b>   <i>Hugging Science (AI-for-Food-Allergies)</i>	Oct 2025 - Nov 2025
<ul style="list-style-type: none"><li>Curated high-quality public datasets for food allergy detection to support AI safety in health domains.</li><li>Engineered an interactive <b>Dataset Explorer</b>, enabling the research community to easily visualize and analyze allergy data distributions.</li></ul>	

## Project Experience

<b>Reasoning Dataset Creation Challenge (Global Competition)</b>	Winner - 1st Place
<ul style="list-style-type: none"><li>Secured <b>1st place globally</b> against 150+ teams in a challenge hosted by HuggingFace, Together.ai, and Bespoke Labs to create synthetic reasoning datasets.</li><li>Designed a novel data synthesis pipeline to generate complex logical reasoning chains, demonstrating superior performance in training Small Language Models (SLMs).</li></ul>	
<b>Kuvera: A Data-Centric Personal Finance LLM &amp; Dataset</b>	HuggingFace Dataset
<ul style="list-style-type: none"><li>Engineered a cost-effective data generation pipeline to fine-tune 8B/14B parameter LLMs (Llama/Mistral) for personal finance.</li><li>Optimized models for efficient single-GPU deployment using LoRA/Quantization, driving <b>31,000+</b> ecosystem downloads (via community GGUF adaptations) and <b>700+</b> base model downloads.</li></ul>	

- PaperStack - A Research Paper Reading Assistant** HuggingFace Space
- Architected a research dashboard that transforms PDF manuscripts into interactive, modular summaries, enhancing information retrieval speed.
  - Implemented a **tiered inference architecture**: enables dynamic switching between a low-latency "Fast Mode" (for rapid screening) and a "Deep Analysis" mode (for comprehensive summarization / context-heavy query resolution), optimizing user cost-performance ratios.

- Causal Analysis of Social Media Signals on Crowdfunding Success** GitHub
- Designed a multi-stage experimental framework to decouple the impact of social engagement metrics from fundamental content features on funding rates.
  - Conducted stepwise ablation studies to quantify the marginal lift of social signals, demonstrating that content features drive success independently of engagement metrics.

## Education

**Indian Institute of Technology, Kharagpur** Aug 2019 – May 2023

*B. Tech, Aerospace Engineering*

### Relevant Coursework:

- Graphical & Generative Modelling for ML
- Dependable & Secure AI-ML
- Linear Algebra for AI and ML
- Theories of Language Comprehension
- Machine Learning Foundations
- Financial Analytics

## Publications

G. Matlin, A. Theerthala, et al. (2025). **FinForge: A Semi-Synthetic Benchmark Generation Framework for Finance.** *Agentic AI in Financial Services, AAAI 2026.*

A. Theerthala (2025). **A Data-Centric Framework for Training Behaviour-Aware Personal Finance Language Models.** *FinNLP Workshop at EMNLP.*

## Certifications

- **The Reasoning Course (Hugging Face)** – *Building reason-capable AI models (2025)*
- **Generative AI with LLMs (DeepLearning.AI)** – *Advanced LLM training & deployment (2025)*
- **Generative AI Nanodegree (Udacity)** – *Comprehensive GenAI pipeline*
- **Machine Learning in Production (Coursera)** – *MLOps & System Design*
- **Deep Learning Specialization (DeepLearning.AI)** – *The core neural networks curriculum*

## Skills

**NLP & LLMs:** HuggingFace (Transformers, TRL, PEFT/LoRA), LangChain, vLLM, Prompt Engineering

**Deep Learning & ML:** PyTorch, TensorFlow, Scikit-learn, OpenCV, NLTK, NumPy, Pandas

**Developer Tools & MLOps:** Docker, Git, Linux

**Visualization:** Matplotlib, Seaborn, Plotly, Gradio/Streamlit